

Notebook

September 26, 2025

```
[1]: #@title Environment & Imports
import math, random, time, os, sys
from dataclasses import dataclass
from typing import Tuple, List, Optional

import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt

# Colab basics
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", device)
if device.type == "cuda":
    print("CUDA:", torch.version.cuda, "| PyTorch:", torch.__version__)
    print("GPU:", torch.cuda.get_device_name(0))
    torch.backends.cudnn.benchmark = True

# Reproducibility
def set_seed(seed: int = 42):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    if device.type == "cuda":
        torch.cuda.manual_seed_all(seed)

set_seed(42)
```

Device: cuda
CUDA: 12.6 | PyTorch: 2.8.0+cu126
GPU: NVIDIA A100-SXM4-40GB

```
[2]: #@title Configuration
from dataclasses import dataclass

@dataclass
class Config:
```

```

H: int = 64
W: int = 64

# Channels: [RGB_LOGITS(3)] + [ALIVE_LOGIT(1)] + [HIDDEN] + [INSTR(3)] + 
↪ [SPARE(1)]
C_VISIBLE: int = 4
C_HIDDEN: int = 16          # a bit more capacity helps stability/growth
C_INSTR: int = 3
C_SPARE: int = 1
C_TOTAL: int = 4 + 16 + 3 + 1

# Liquid gate (closed-form) hyperparams
DELTA_MIN: float = 0.02

# Training
BATCH: int = 16
LR: float = 2e-3
WD: float = 1e-6
MAX_ITERS_A: int = 1500
MAX_ITERS_B: int = 800
ROLLOUT_T_MIN: int = 48
ROLLOUT_T_MAX: int = 96
MID_LOSS_W: float = 0.30    # weight on mid-step loss

# Damage
DAMAGE_FRAC_MIN: float = 0.20
DAMAGE_FRAC_MAX: float = 0.35
DAMAGE_PUSH: float = 6.0

# Loss weights
LAMBDA_MSE: float = 1.0
LAMBDA_TV: float = 0.02
LAMBDA_ALIVE_MEAN: float = 1e-4
LAMBDA_H_NORM: float = 0.001
FG_WEIGHT: float = 8.0      # foreground up-weight for thin rings
ALIVE_BCE_W: float = 0.7    # supervision on alive logit
ALIVE_POS_WEIGHT: float = 6.0 # BCE pos_weight for thin foreground

# Local update mask (from alive prob)
ALIVE_THRESH: float = 0.10
ALIVE_KAPPA: float = 10.0   # steepness of sigmoid around thresh
LOGIT_LEAK: float = 0.01    # tiny decay on RGB logits per step

# Geometry: fix for Day-1 to learn a concrete rule first
FIXED_GEOMETRY: bool = True
RING_R: float = 0.55
RING_TH: float = 0.10

```

```

DIAMOND_R: float = 0.65
DIAMOND_TH: float = 0.08
BILOBE_R: float = 0.50
BILOBE_D: float = 0.60
BILOBE_TH: float = 0.10

VIZ_STEPS: int = 64

cfg = Config()
print(cfg)

```

```

Config(H=64, W=64, C_VISIBLE=4, C_HIDDEN=16, C_INSTR=3, C_SPARE=1, C_TOTAL=24,
DELTA_MIN=0.02, BATCH=16, LR=0.002, WD=1e-06, MAX_ITERS_A=1500, MAX_ITERS_B=800,
ROLLOUT_T_MIN=48, ROLLOUT_T_MAX=96, MID_LOSS_W=0.3, DAMAGE_FRAC_MIN=0.2,
DAMAGE_FRAC_MAX=0.35, DAMAGE_PUSH=6.0, LAMBDA_MSE=1.0, LAMBDA_TV=0.02,
LAMBDA_ALIVE_MEAN=0.0001, LAMBDA_H_NORM=0.001, FG_WEIGHT=8.0, ALIVE_BCE_W=0.7,
ALIVE_POS_WEIGHT=6.0, ALIVE_THRESH=0.1, ALIVE_KAPPA=10.0, LOGIT_LEAK=0.01,
FIXED_GEOMETRY=True, RING_R=0.55, RING_TH=0.1, DIAMOND_R=0.65, DIAMOND_TH=0.08,
BILOBE_R=0.5, BILOBE_D=0.6, BILOBE_TH=0.1, VIZ_STEPS=64)

```

[3]: *#@title Utility: Targets, Seed, Losses, Damage*

```

import torch.nn.functional as F

def make_coords(H, W, device=device):
    ys = torch.linspace(-1.0, 1.0, H, device=device).view(H, 1).expand(H, W)
    xs = torch.linspace(-1.0, 1.0, W, device=device).view(1, W).expand(H, W)
    return xs, ys

XS, YS = make_coords(cfg.H, cfg.W)

def target_ring(H: int, W: int, thickness: float, center, r: float):
    cx, cy = center
    dist = torch.sqrt((XS - cx)**2 + (YS - cy)**2)
    mask = (dist > (r - thickness)) & (dist < (r + thickness))
    return mask.float()

def target_diamond(H: int, W: int, thickness: float, center, r: float):
    cx, cy = center
    dist1 = (torch.abs(XS - cx) + torch.abs(YS - cy))
    mask = (dist1 > (r - thickness)) & (dist1 < (r + thickness))
    return mask.float()

def target_bilobe(H: int, W: int, thickness: float, center, r: float, d: float):
    cx, cy = center
    x1, y1 = cx - d/2, cy
    x2, y2 = cx + d/2, cy

```

```

    dist = torch.sqrt((XS - x1)**2 + (YS - y1)**2) + torch.sqrt((XS - x2)**2 +
↪(YS - y2)**2)
    mask = (dist < (2*r + thickness)) & (dist > (2*r - thickness))
    return mask.float()

def _pix_to_norm(cx_pix: int, cy_pix: int, H: int, W: int):
    cx = 2.0 * (cx_pix / (W - 1)) - 1.0
    cy = 2.0 * (cy_pix / (H - 1)) - 1.0
    return cx, cy

def init_state(batch: int, cfg: Config, instr_onehot: torch.Tensor):
    """
    Returns:
        state: (B, C_TOTAL, H, W) -> visible are LOGITS
        centers: (B, 2) normalized centers
    """
    B, H, W = batch, cfg.H, cfg.W
    state = torch.zeros((B, cfg.C_TOTAL, H, W), device=device)

    # Initialize background: dark & dead
    state[:, :3, :, :] = -4.0      # RGB logits ~ 0.018 prob
    state[:, 3:4, :, :] = -6.0    # ALIVE logit ~ 0.002 prob

    # Instruction broadcast (filled after we sample class)
    instr_grid = instr_onehot.view(B, cfg.C_INSTR, 1, 1).expand(B, cfg.C_INSTR,
↪H, W)
    state[:, cfg.C_VISIBLE + cfg.C_HIDDEN : cfg.C_VISIBLE + cfg.C_HIDDEN + cfg.
↪C_INSTR] = instr_grid

    centers = []
    for b in range(B):
        cx_pix = random.randint(8, W - 9)
        cy_pix = random.randint(8, H - 9)
        # Seed: alive on, RGB slightly positive to kick-start
        state[b, 3, cy_pix-1:cy_pix+2, cx_pix-1:cx_pix+2] = 4.0
        state[b, :3, cy_pix-1:cy_pix+2, cx_pix-1:cx_pix+2] = 0.5
        state[b, -1, cy_pix-1:cy_pix+2, cx_pix-1:cx_pix+2] = 1.0 # spare seed
↪mark
        centers.append(_pix_to_norm(cx_pix, cy_pix, H, W))

    centers = torch.tensor(centers, dtype=torch.float32, device=device)
    return state, centers

def sample_target_and_instr(batch: int, H: int, W: int, centers: torch.Tensor,
↪device=device):
    imgs, instrs = [], []
    for i in range(batch):

```

```

        cx, cy = centers[i].tolist()
        k = random.randint(0, 2)
        if k == 0:
            if cfg.FIXED_GEOMETRY:
                g = target_ring(H, W, thickness=cfg.RING_TH, center=(cx, cy),
↪r=cfg.RING_R)
            else:
                r = random.uniform(0.35, 0.60); th = random.uniform(0.08, 0.12)
                g = target_ring(H, W, thickness=th, center=(cx, cy), r=r)
        elif k == 1:
            if cfg.FIXED_GEOMETRY:
                g = target_diamond(H, W, thickness=cfg.DIAMOND_TH, center=(cx,
↪cy), r=cfg.DIAMOND_R)
            else:
                r = random.uniform(0.50, 0.75); th = random.uniform(0.07, 0.12)
                g = target_diamond(H, W, thickness=th, center=(cx, cy), r=r)
        else:
            if cfg.FIXED_GEOMETRY:
                g = target_bilobe(H, W, thickness=cfg.BILOBE_TH, center=(cx,
↪cy), r=cfg.BILOBE_R, d=cfg.BILOBE_D)
            else:
                r = random.uniform(0.45, 0.60); d = random.uniform(0.5, 0.7);
↪th = random.uniform(0.08, 0.12)
                g = target_bilobe(H, W, thickness=th, center=(cx, cy), r=r, d=d)

        rgb = torch.stack([g, g, g], dim=0).to(device)
        imgs.append(rgb)
        instrs.append(torch.tensor([1,0,0] if k==0 else [0,1,0] if k==1 else
↪[0,0,1],
                                device=device, dtype=torch.float32))
        targets = torch.stack(imgs, dim=0) # (B, 3, H, W) in {0,1}
        instr = torch.stack(instrs, dim=0) # (B, 3)
        return targets, instr

# Losses -----

def weighted_mse(pred, target, fg_weight: float):
    fg = target[:, :1, :, :]
    w = 1.0 + fg_weight * fg
    return ((pred - target)**2 * w).mean()

def tv_loss(img):
    dx = img[:, :, 1:, :] - img[:, :, :-1, :]
    dy = img[:, :, :, 1:] - img[:, :, :, :-1]
    return (dx.abs().mean() + dy.abs().mean())

def hidden_norm_penalty(h):

```

```

        return (h**2).mean()

# Damage: push logits strongly negative in region
def apply_random_damage(state, cfg: Config, frac_range=(0.2, 0.35)):
    B, _, H, W = state.shape
    fmin, fmax = frac_range
    for b in range(B):
        fw = random.uniform(fmin, fmax); fh = random.uniform(fmin, fmax)
        ww = max(1, int(W * fw)); hh = max(1, int(H * fh))
        x0 = random.randint(0, W - ww); y0 = random.randint(0, H - hh)
        state[b, :3, y0:y0+hh, x0:x0+ww] -= cfg.DAMAGE_PUSH
        state[b, 3:4, y0:y0+hh, x0:x0+ww] -= cfg.DAMAGE_PUSH
    return state

```

```

[4]: #@title Liquid Cell + Neighborhood
class Neighborhood(nn.Module):
    def __init__(self, C: int, k: int = 3, padding: int = 1):
        super().__init__()
        self.k = k
        self.padding = padding
        self.C = C

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        B, C, H, W = x.shape
        patches = F.unfold(x, kernel_size=self.k, padding=self.padding,
↪stride=1)
        return patches.view(B, C*self.k*self.k, H, W)

class LiquidCell(nn.Module):
    """
    Visible channels are LOGITS: [rgb_logits(3), alive_logit(1)].
    We:
    - compute alive prob -> local update mask via softened 3x3 pooling
    - update hidden & visible only under that mask
    - apply tiny leak on RGB logits to avoid saturation
    """
    def __init__(self, cfg: Config):
        super().__init__()
        self.cfg = cfg
        self.neigh = Neighborhood(cfg.C_TOTAL, k=3, padding=1)

        Cn = cfg.C_TOTAL * 9
        Hh = cfg.C_HIDDEN
        Cv = cfg.C_VISIBLE

        self.Wx = nn.Conv2d(Cn, Hh, kernel_size=1, bias=True)
        self.Wh = nn.Conv2d(Hh, Hh, kernel_size=1, bias=False)
        self.Wg = nn.Conv2d(Cn + Hh, Hh, kernel_size=1, bias=True)

```

```

self.Wout = nn.Conv2d(Hh + Cn, Cv, kernel_size=1, bias=True)

for m in [self.Wx, self.Wh, self.Wg, self.Wout]:
    nn.init.xavier_uniform_(m.weight, gain=0.7)
    if m.bias is not None: nn.init.zeros_(m.bias)
self.Wout.weight.data *= 0.25

self.delta_min = cfg.DELTA_MIN
self.kappa = cfg.ALIVE_KAPPA
self.thresh = cfg.ALIVE_THRESH
self.leak = cfg.LOGIT_LEAK

def forward(self, state: torch.Tensor) -> torch.Tensor:
    B, C, H, W = state.shape
    Cv, Hh = self.cfg.C_VISIBLE, self.cfg.C_HIDDEN

    neigh = self.neigh(state)          # (B, 9*C, H, W)
    logits = state[:, :Cv]              # (B, 4, H, W)
    rgb_logits = logits[:, :3]
    alive_logits = logits[:, 3:4]
    h = state[:, Cv:Cv+Hh]
    rest = state[:, Cv+Hh:]

    # Alive-based local update mask
    alive_prob = torch.sigmoid(alive_logits)
    avg_alive = F.avg_pool2d(alive_prob, kernel_size=3, stride=1, padding=1)
    update_mask = torch.sigmoid(self.kappa * (avg_alive - self.thresh)) #␣
    ↪(B,1,H,W)

    # Hidden update (liquid gate) under mask
    h_tilde = torch.tanh(self.Wx(neigh) + self.Wh(h))
    delta_raw = torch.sigmoid(self.Wg(torch.cat([neigh, h], dim=1)))
    delta = delta_raw * (1 - 2*self.delta_min) + self.delta_min
    h_next = h + update_mask * delta * (h_tilde - h)

    # Visible logits update under mask (+ tiny leak on RGB)
    out_in = torch.cat([h_next, neigh], dim=1)
    dlogits = self.Wout(out_in)
    rgb_logits_next = (1 - self.leak) * rgb_logits + update_mask *␣
    ↪dlogits[:, :3]
    alive_logits_next = alive_logits + update_mask * dlogits[:, 3:4]

    new_visible = torch.cat([rgb_logits_next, alive_logits_next], dim=1)
    new_state = torch.cat([new_visible, h_next, rest], dim=1)
    return new_state

```

```
[5]: #@title Model Wrapper & Training Rollout
class LiquidMorphCA(nn.Module):
    def __init__(self, cfg: Config):
        super().__init__()
        self.cfg = cfg
        self.cell = LiquidCell(cfg)

    @torch.no_grad()
    def step_no_grad(self, state: torch.Tensor, steps: int = 1):
        for _ in range(steps): state = self.cell(state)
        return state

    def step(self, state: torch.Tensor, steps: int = 1):
        for _ in range(steps): state = self.cell(state)
        return state

def compute_losses(state: torch.Tensor, target_rgb: torch.Tensor, cfg: Config):
    Cv, Hh = cfg.C_VISIBLE, cfg.C_HIDDEN
    rgb_logits = state[:, :3]
    alive_logits = state[:, 3:4]
    hidden = state[:, Cv:Cv+Hh]

    pred_rgb = torch.sigmoid(rgb_logits)
    alive_prob = torch.sigmoid(alive_logits)
    target_mask = target_rgb[:, :1] # ring/edge mask in {0,1}

    # RGB reconstruction (foreground-weighted) + smoothness
    loss_rgb = weighted_mse(pred_rgb, target_rgb, cfg.FG_WEIGHT)
    loss_tv = tv_loss(pred_rgb)

    # Alive supervision (BCE with class imbalance)
    pos_w = torch.tensor(cfg.ALIVE_POS_WEIGHT, device=state.device)
    loss_alive_sup = F.binary_cross_entropy_with_logits(alive_logits,
    ↪target_mask, pos_weight=pos_w)

    # Regularizers
    loss_alive_mean = alive_prob.mean()
    loss_hnorm = hidden_norm_penalty(hidden)

    total = (cfg.LAMBDA_MSE * loss_rgb
            + cfg.LAMBDA_TV * loss_tv
            + cfg.ALIVE_BCE_W * loss_alive_sup
            + cfg.LAMBDA_ALIVE_MEAN * loss_alive_mean
            + cfg.LAMBDA_H_NORM * loss_hnorm)

    stats = {
        "loss_total": float(total.detach()),
```



```

        "loss_rgb": float(loss_rgb.detach()),
        "loss_tv": float(loss_tv.detach()),
        "loss_alive_sup": float(loss_alive_sup.detach()),
        "alive_mean": float(loss_alive_mean.detach()),
        "hnorm": float(loss_hnorm.detach()),
    }
    return total, stats

def rollout_training(model: LiquidMorphCA,
                    cfg: Config,
                    stage: str = "A",
                    optimizer: Optional[torch.optim.Optimizer] = None):
    model.train()
    B, H, W = cfg.BATCH, cfg.H, cfg.W

    dummy_instr = torch.zeros((B, cfg.C_INSTR), device=device)
    state, centers = init_state(B, cfg, dummy_instr)
    targets, instr = sample_target_and_instr(B, H, W, centers, device=device)

    # write instruction channels
    instr_grid = instr.view(B, cfg.C_INSTR, 1, 1).expand(B, cfg.C_INSTR, H, W)
    state[:, cfg.C_VISIBLE + cfg.C_HIDDEN : cfg.C_VISIBLE + cfg.C_HIDDEN + cfg.
↪C_INSTR] = instr_grid

    T = random.randint(cfg.ROLLOUT_T_MIN, cfg.ROLLOUT_T_MAX)
    damage_step = random.randint(T//4, (2*T)//3) if stage == "B" else None
    mid_t = int(0.75 * T)
    state_mid = None

    for t in range(T):
        state = model.step(state, steps=1)
        if t == mid_t:
            state_mid = state
        if stage == "B" and t == damage_step:
            state = apply_random_damage(state, cfg, (cfg.DAMAGE_FRAC_MIN, cfg.
↪DAMAGE_FRAC_MAX))

    loss_end, stats_end = compute_losses(state, targets, cfg)
    if state_mid is not None:
        loss_mid, _ = compute_losses(state_mid, targets, cfg)
        total = (1 - cfg.MID_LOSS_W) * loss_end + cfg.MID_LOSS_W * loss_mid
        stats_end["loss_mid"] = float(loss_mid.detach())
    else:
        total = loss_end

    if optimizer is not None:
        optimizer.zero_grad(set_to_none=True)

```

```

total.backward()
torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
optimizer.step()

stats_end["T"] = T
stats_end["loss_total"] = float(total.detach())
return stats_end

```

```

[6]: #@title Visualization Helpers (UPDATED)
from IPython.display import clear_output

def _rgb_from_logits(rgb_logits: torch.Tensor):
    rgb = torch.sigmoid(rgb_logits).clamp(0,1)
    return rgb.permute(1,2,0).detach().cpu().numpy()

@torch.no_grad()
def render_growth(model: LiquidMorphCA,
                  cfg: Config,
                  pattern_idx: int = 0,
                  steps: int = 64,
                  damage_at: Optional[int] = None,
                  damage_frac=(0.25, 0.3),
                  show_alive: bool = False,
                  seed=None):
    if seed is not None:
        set_seed(seed)

    B, H, W = 1, cfg.H, cfg.W
    instr = torch.tensor([[1,0,0] if pattern_idx==0 else [0,1,0] if
↪pattern_idx==1 else [0,0,1]],
                        dtype=torch.float32, device=device)
    state, centers = init_state(B, cfg, instr)
    cx, cy = centers[0].tolist()

    if pattern_idx == 0:
        g = target_ring(H, W, thickness=cfg.RING_TH, center=(cx, cy), r=cfg.
↪RING_R)
    elif pattern_idx == 1:
        g = target_diamond(H, W, thickness=cfg.DIAMOND_TH, center=(cx, cy),
↪r=cfg.DIAMOND_R)
    else:
        g = target_bilobe(H, W, thickness=cfg.BILOBE_TH, center=(cx, cy), r=cfg.
↪BILOBE_R, d=cfg.BILOBE_D)
    target = torch.stack([g,g,g], dim=0).unsqueeze(0).to(device)

    instr_grid = instr.view(B, cfg.C_INSTR, 1, 1).expand(B, cfg.C_INSTR, H, W)

```

```

state[:, cfg.C_VISIBLE + cfg.C_HIDDEN : cfg.C_VISIBLE + cfg.C_HIDDEN + cfg.
↪C_INSTR] = instr_grid

frames = []
for t in range(steps):
    state = model.step_no_grad(state, steps=1)
    if damage_at is not None and t == damage_at:
        state = apply_random_damage(state, cfg, damage_frac)

    img = _rgb_from_logits(state[0, :3])
    if show_alive:
        alive = torch.sigmoid(state[0, 3:4]).repeat(3,1,1).permute(1,2,0).
↪cpu().numpy()
        img = (0.8*img + 0.2*np.stack([alive[:, :, 0], np.zeros_like(alive[:, :,
↪, 0]), np.zeros_like(alive[:, :, 0])], axis=-1))
        frames.append(img)

    fig, axes = plt.subplots(1, 3, figsize=(10,3))
    axes[0].imshow(frames[0]); axes[0].set_title("Step 0"); axes[0].
↪axis('off')
    axes[1].imshow(frames[steps//2]); axes[1].set_title(f"Step {steps//2}");
↪axes[1].axis('off')
    axes[2].imshow(frames[-1]); axes[2].set_title(f"Step {steps-1}");
↪axes[2].axis('off')
    plt.show()

return frames, target[0]

```

```

[7]: #@title Instantiate Model & Optimizer
model = LiquidMorphCA(cfg).to(device)
optimizer = torch.optim.AdamW(model.parameters(), lr=cfg.LR, weight_decay=cfg.
↪WD)
print("Parameters:", sum(p.numel() for p in model.parameters())/1e6, "M")

```

Parameters: 0.008388 M

```

[8]: #@title Stage A Training (Growth)
loss_log_A = []
start = time.time()
for it in range(1, cfg.MAX_ITERS_A + 1):
    stats = rollout_training(model, cfg, stage="A", optimizer=optimizer)
    loss_log_A.append(stats["loss_total"])

    if it % 100 == 0 or it == 1:
        clear_output(wait=True)
        dt = time.time() - start

```

```

# First iteration: print all available stats to see what we have
if it == 1:
    print("Available stats:", stats.keys())

# Simple version - just print what we have
print(f"[Stage A] iter={it}/{cfg.MAX_ITERS_A} | T={stats['T']} | "
      f"loss={stats['loss_total']:.4f} | {dt:.1f}s")

# Optional: print all loss components
for key, val in stats.items():
    if key.startswith('loss_') and key != 'loss_total':
        print(f"  {key}: {val:.6f}")

with torch.no_grad():
    frames, tgt = render_growth(model, cfg, pattern_idx=random.
    ↪ randint(0,2),
                                steps=cfg.VIZ_STEPS, seed=it)

    plt.figure(figsize=(4,4))
    plt.imshow(_rgb_from_logits(tgt)); plt.title("Target Example"); plt.
    ↪ axis('off')
    plt.show()

```

```

[Stage A] iter=1500/1500 | T=67 | loss=0.0409 | 285.2s
loss_rgb: 0.013797
loss_tv: 0.055682
loss_alive_sup: 0.032711
loss_mid: 0.046817

```

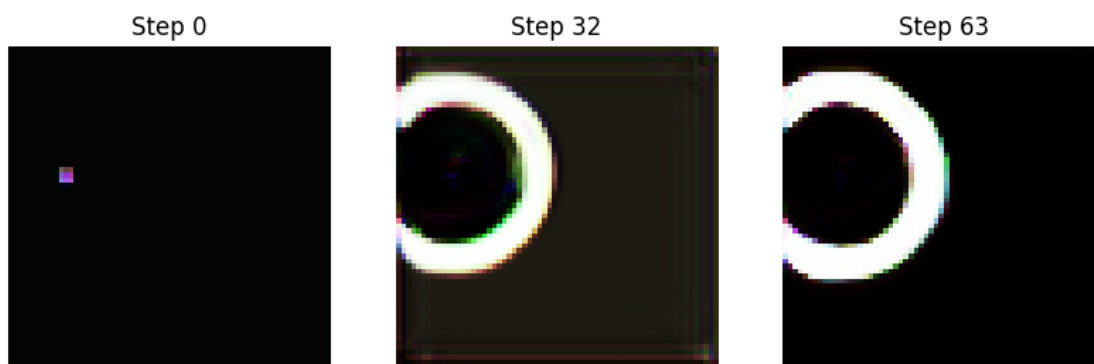


Target Example



```
[9]: #@title Visual Check: Growth (No Damage)
with torch.no_grad():
    for k in range(3):
        print(["Ring", "Diamond", "Bi-lobe"][k])
        _ = render_growth(model, cfg, pattern_idx=k, steps=cfg.VIZ_STEPS,
        ↪damage_at=None, seed=123+k)
```

Ring



Diamond



Bi-lobe



```
[10]: #@title Stage B Training (Regeneration)
loss_log_B = []
start = time.time()
for it in range(1, cfg.MAX_ITERS_B + 1):
    stats = rollout_training(model, cfg, stage="B", optimizer=optimizer)
    loss_log_B.append(stats["loss_total"])

    if it % 100 == 0 or it == 1:
        clear_output(wait=True)
        dt = time.time() - start

        # First iteration: print all available stats to see what we have
        if it == 1:
            print("Available stats:", stats.keys())

        # Simple version - just print what we have
        print(f"[Stage B] iter={it}/{cfg.MAX_ITERS_B} | T={stats['T']} | "
              f"loss={stats['loss_total']:.4f} | {dt:.1f}s")
```

```

# Optional: print all loss components
for key, val in stats.items():
    if key.startswith('loss_') and key != 'loss_total':
        print(f" {key}: {val:.6f}")

with torch.no_grad():
    frames, _ = render_growth(model, cfg, pattern_idx=random.
↳ randint(0,2),
                                steps=cfg.VIZ_STEPS, damage_at=cfg.
↳ VIZ_STEPS//3,
                                damage_frac=(cfg.DAMAGE_FRAC_MIN, cfg.
↳ DAMAGE_FRAC_MAX),
                                seed=777+it)

```

[Stage B] iter=800/800 | T=81 | loss=0.0615 | 154.0s
 loss_rgb: 0.019991
 loss_tv: 0.062161
 loss_alive_sup: 0.062732
 loss_mid: 0.051347

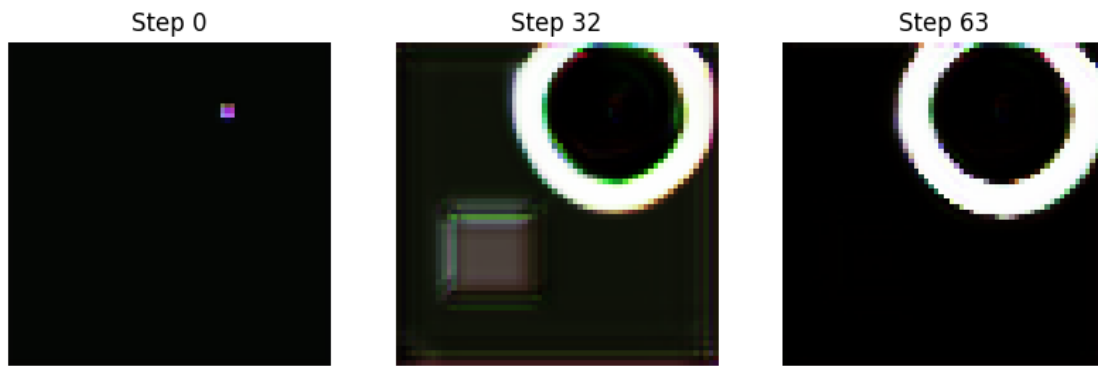


```

[11]: #@title Visual Check: Regeneration
with torch.no_grad():
    for k in range(3):
        print(["Ring", "Diamond", "Bi-lobe"][k], "with mid-rollout damage")
        _ = render_growth(model, cfg, pattern_idx=k, steps=cfg.VIZ_STEPS,
                            damage_at=cfg.VIZ_STEPS//3, damage_frac=(0.25, 0.3),
↳ seed=2025+k)

```

Ring with mid-rollout damage



Diamond with mid-rollout damage



Bi-lobe with mid-rollout damage



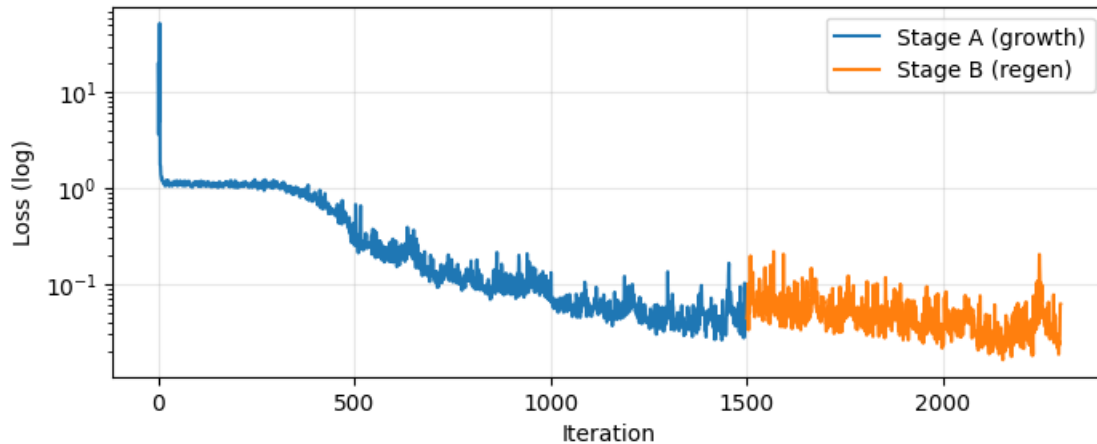
```
[12]: #@title Loss Curves
plt.figure(figsize=(8,3))
```



```

if len(loss_log_A):
    plt.plot(loss_log_A, label="Stage A (growth)")
if len(loss_log_B):
    plt.plot(range(len(loss_log_A), len(loss_log_A)+len(loss_log_B)),
             loss_log_B, label="Stage B (regen)")
plt.yscale('log')
plt.xlabel("Iteration"); plt.ylabel("Loss (log)"); plt.legend(); plt.grid(True,
    alpha=0.3)
plt.show()

```



```

[13]: #@title (Optional) Ablation: Fixed Gate
class FixedGateCell(LiquidCell):
    def __init__(self, cfg: Config, alpha: float = 0.25):
        super().__init__(cfg)
        self.alpha = alpha
    def forward(self, state):
        B, C, H, W = state.shape
        Cv, Hh = self.cfg.C_VISIBLE, self.cfg.C_HIDDEN
        neigh = self.neigh(state)
        h = state[:, Cv:Cv+Hh]
        y = state[:, :Cv]
        rest = state[:, Cv+Hh:]

        h_tilde = torch.tanh(self.Wx(neigh) + self.Wh(h))
        # Fixed gate (no input-dependent adaptation)
        delta = torch.full_like(h, self.alpha)
        h_next = h + delta * (h_tilde - h)

        out_in = torch.cat([h_next, neigh], dim=1)
        dy = self.Wout(out_in)

```

```

        y_next = (y + dy).clamp(0,1)
        return torch.cat([y_next, h_next, rest], dim=1)

# Quick comparison (no training; just copy weights from trained model and
↳ replace cell)
def clone_to_fixed_gate(trained_model: LiquidMorphCA, alpha=0.25):
    m = LiquidMorphCA(cfg).to(device)
    m.load_state_dict(trained_model.state_dict(), strict=True)
    fg = FixedGateCell(cfg, alpha=alpha).to(device)
    # Copy learned weights into the fixed-gate cell
    fg.Wx.load_state_dict(trained_model.cell.Wx.state_dict())
    fg.Wh.load_state_dict(trained_model.cell.Wh.state_dict())
    fg.Wout.load_state_dict(trained_model.cell.Wout.state_dict())
    m.cell = fg
    return m

fixed_model = clone_to_fixed_gate(model, alpha=0.25)
with torch.no_grad():
    print("Fixed-gate behavior (compare visually to liquid):")
    _ = render_growth(fixed_model, cfg, pattern_idx=0, steps=cfg.VIZ_STEPS,
↳ damage_at=cfg.VIZ_STEPS//3, seed=909)

```

Fixed-gate behavior (compare visually to liquid):



[14]: *#@title (a) Parametrized size/rotation targets*

```

import math

# Ranges for generalization (you can tweak live)
cfg.ROT_MIN, cfg.ROT_MAX = -math.pi/4, math.pi/4      # +/- 45 degrees
cfg.SCALE_MIN, cfg.SCALE_MAX = 0.75, 1.25             # 75%..125% of base size

def rotate_coords(XS, YS, center, theta):

```

```

"""Rotate coordinates around center by theta (radians)."""
cx, cy = center
x0, y0 = XS - cx, YS - cy
ct, st = math.cos(theta), math.sin(theta)
xr = x0 * ct - y0 * st
yr = x0 * st + y0 * ct
return xr + cx, yr + cy

def target_ring_param(H, W, center, base_r, thickness, scale=1.0, theta=0.0):
    # rotation has no effect for ring; keep interface consistent
    r = base_r * scale
    dist = torch.sqrt((XS - center[0])**2 + (YS - center[1])**2)
    mask = (dist > (r - thickness)) & (dist < (r + thickness))
    return mask.float()

def target_diamond_param(H, W, center, base_r, thickness, scale=1.0, theta=0.0):
    # diamond defined by L1 radius; rotation via rotated coords
    r = base_r * scale
    Xr, Yr = rotate_coords(XS, YS, center, theta)
    dist1 = torch.abs(Xr - center[0]) + torch.abs(Yr - center[1])
    mask = (dist1 > (r - thickness)) & (dist1 < (r + thickness))
    return mask.float()

def target_bilobe_param(H, W, center, base_r, base_d, thickness, scale=1.0,
    ↪theta=0.0):
    # two foci separated by d; apply rotation by rotating the *offset axis*
    r = base_r * scale
    d = base_d * scale
    # axis endpoints before rotation
    x1, y1 = center[0] - d/2, center[1]
    x2, y2 = center[0] + d/2, center[1]
    # rotate endpoints around center
    ct, st = math.cos(theta), math.sin(theta)
    def rot(x, y):
        dx, dy = x - center[0], y - center[1]
        xr = dx * ct - dy * st + center[0]
        yr = dx * st + dy * ct + center[1]
        return xr, yr
    x1r, y1r = rot(x1, y1); x2r, y2r = rot(x2, y2)
    dist = torch.sqrt((XS - x1r)**2 + (YS - y1r)**2) + torch.sqrt((XS - x2r)**2
    ↪+ (YS - y2r)**2)
    mask = (dist < (2*r + thickness)) & (dist > (2*r - thickness))
    return mask.float()

def make_phi_field(shape_id, center, scale, theta):
    """
    A steering potential phi(x,y) whose small values lie ON the desired contour.

```

```

We store phi in the SPARE channel so the rule can read 'where to grow'.
"""

if shape_id == 0: # ring
    r = cfg.RING_R * scale
    dist = torch.sqrt((XS - center[0])**2 + (YS - center[1])**2)
    phi = torch.abs(dist - r) / (r + 1e-6)
elif shape_id == 1: # diamond
    Xr, Yr = rotate_coords(XS, YS, center, theta)
    r = cfg.DIAMOND_R * scale
    dist1 = torch.abs(Xr - center[0]) + torch.abs(Yr - center[1])
    phi = torch.abs(dist1 - r) / (r + 1e-6)
else: # bilobe
    r = cfg.BILOBE_R * scale
    d = cfg.BILOBE_D * scale
    # rotated endpoints
    ct, st = math.cos(theta), math.sin(theta)
    def rot(x, y):
        dx, dy = x - center[0], y - center[1]
        xr = dx * ct - dy * st + center[0]
        yr = dx * st + dy * ct + center[1]
        return xr, yr
    x1, y1 = center[0] - d/2, center[1]
    x2, y2 = center[0] + d/2, center[1]
    x1r, y1r = rot(x1, y1); x2r, y2r = rot(x2, y2)
    dist = torch.sqrt((XS - x1r)**2 + (YS - y1r)**2) + torch.sqrt((XS -
↪x2r)**2 + (YS - y2r)**2)
    phi = torch.abs(dist - 2*r) / (2*r + 1e-6)
    # normalize to [0,1]
    phi = (phi / (phi.max() + 1e-6)).clamp(0, 1)
    return phi

def sample_param_targets(batch, H, W, centers, device=device):
    """Random size/rotation per sample; also returns phi fields."""
    imgs, instrs, phis = [], [], []
    for i in range(batch):
        cx, cy = centers[i].tolist()
        k = random.randint(0, 2) # shape id
        theta = random.uniform(cfg.ROT_MIN, cfg.ROT_MAX)
        scale = random.uniform(cfg.SCALE_MIN, cfg.SCALE_MAX)

        if k == 0:
            g = target_ring_param(H, W, (cx, cy), cfg.RING_R, cfg.RING_TH,
↪scale, theta)
        elif k == 1:
            g = target_diamond_param(H, W, (cx, cy), cfg.DIAMOND_R, cfg.
↪DIAMOND_TH, scale, theta)
        else:

```

```

        g = target_bilobe_param(H, W, (cx, cy), cfg.BILOBE_R, cfg.BILOBE_D,
↪cfg.BILOBE_TH, scale, theta)

        phi = make_phi_field(k, (cx, cy), scale, theta)
        rgb = torch.stack([g, g, g], dim=0).to(device)
        imgs.append(rgb)
        instrs.append(torch.tensor([1,0,0] if k==0 else [0,1,0] if k==1 else
↪[0,0,1],
                                device=device, dtype=torch.float32))
        phis.append(phi.to(device))

    targets = torch.stack(imgs, dim=0)           # (B, 3, H, W)
    instr  = torch.stack(instrs, dim=0)          # (B, 3)
    phi    = torch.stack(phis, dim=0)            # (B, H, W) in [0,1]
    return targets, instr, phi

```

[15]: *#@title (a) Training stage A-gen: size/rotation generalization (uses field)*

```

def rollout_training_param(model: LiquidMorphCA,
                           cfg: Config,
                           optimizer: torch.optim.Optimizer,
                           iters: int = 800):

    loss_log = []
    for it in range(1, iters+1):
        B, H, W = cfg.BATCH, cfg.H, cfg.W
        dummy_instr = torch.zeros((B, cfg.C_INSTR), device=device)
        state, centers = init_state(B, cfg, dummy_instr)

        # new parametrized targets + phi field
        targets, instr, phi = sample_param_targets(B, H, W, centers,
↪device=device)

        # write instruction
        instr_grid = instr.view(B, cfg.C_INSTR, 1, 1).expand(B, cfg.C_INSTR, H,
↪W)

        state[:, cfg.C_VISIBLE + cfg.C_HIDDEN : cfg.C_VISIBLE + cfg.C_HIDDEN +
↪cfg.C_INSTR] = instr_grid

        # write phi into SPARE channel (overrides seed mark)
        state[:, -1] = phi

        # rollout
        T = random.randint(cfg.ROLLOUT_T_MIN, cfg.ROLLOUT_T_MAX)
        mid_t = int(0.75 * T)
        state_mid = None
        for t in range(T):
            state = model.step(state, steps=1)
            if t == mid_t: state_mid = state

```

```

    loss_end, stats_end = compute_losses(state, targets, cfg)
    if state_mid is not None:
        loss_mid, _ = compute_losses(state_mid, targets, cfg)
        total = (1 - cfg.MID_LOSS_W) * loss_end + cfg.MID_LOSS_W * loss_mid
    else:
        total = loss_end

    optimizer.zero_grad(set_to_none=True)
    total.backward()
    torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
    optimizer.step()

    loss_log.append(float(total.detach()))
    if it % 100 == 0 or it == 1:
        print(f"[A-gen] it={it}/{iters} | loss={float(total.detach()):.4f} |
↪| "
                f"loss_rgb={stats_end['loss_rgb']:.4f} |
↪alive_sup={stats_end['loss_alive_sup']:.4f}")
    return loss_log

```

[16]: *#@title (a) Visualization: parametrized size/rotation*

```

@torch.no_grad()
def render_growth_param(model: LiquidMorphCA,
                        cfg: Config,
                        shape_id: int = 0,
                        scale: float = 1.0,
                        theta: float = 0.0,
                        steps: int = 64,
                        seed=None):
    if seed is not None: set_seed(seed)
    B, H, W = 1, cfg.H, cfg.W
    instr = torch.tensor([[1,0,0] if shape_id==0 else [0,1,0] if shape_id==1
↪else [0,0,1]],
                        dtype=torch.float32, device=device)
    state, centers = init_state(B, cfg, instr)
    cx, cy = centers[0].tolist()

    # target + phi
    if shape_id == 0:
        g = target_ring_param(H, W, (cx,cy), cfg.RING_R, cfg.RING_TH, scale,
↪theta)
    elif shape_id == 1:
        g = target_diamond_param(H, W, (cx,cy), cfg.DIAMOND_R, cfg.DIAMOND_TH,
↪scale, theta)
    else:

```

```

        g = target_bilobe_param(H, W, (cx,cy), cfg.BILOBE_R, cfg.BILOBE_D, cfg.
↪BILOBE_TH, scale, theta)
        target = torch.stack([g,g,g], dim=0).unsqueeze(0).to(device)
        phi = make_phi_field(shape_id, (cx,cy), scale, theta)
        # write instr + phi
        instr_grid = instr.view(B, cfg.C_INSTR, 1, 1).expand(B, cfg.C_INSTR, H, W)
        state[:, cfg.C_VISIBLE + cfg.C_HIDDEN : cfg.C_VISIBLE + cfg.C_HIDDEN + cfg.
↪C_INSTR] = instr_grid
        state[:, -1] = phi

        frames = []
        for t in range(steps):
            state = model.step_no_grad(state, steps=1)
            frames.append(torch.sigmoid(state[0,:3]).permute(1,2,0).cpu().numpy())

        fig, axes = plt.subplots(1, 3, figsize=(10,3))
        axes[0].imshow(frames[0]); axes[0].set_title("Step 0"); axes[0].
↪axis('off')
        axes[1].imshow(frames[steps//2]); axes[1].set_title(f"Step {steps//2}");
↪axes[1].axis('off')
        axes[2].imshow(frames[-1]); axes[2].set_title(f"Step {steps-1}");
↪axes[2].axis('off')
        plt.show()

# Example (you can tweak)
# _ = rollout_training_param(model, cfg, optimizer, iters=400)
# render_growth_param(model, cfg, shape_id=1, scale=1.15, theta=0.35, steps=64,
↪seed=123)

```

[17]: #@title (b) Chemical environment and diffusion update

```

from dataclasses import dataclass

@dataclass
class ChemParams:
    D: float = 0.22          # diffusion
    decay: float = 0.01      # linear decay
    n_sources: int = 2       # number of Gaussian sources
    strength_min: float = 0.8
    strength_max: float = 1.2
    sigma_min: float = 0.08  # as fraction of [-1,1] domain
    sigma_max: float = 0.18
    threshold: float = 0.35  # binarization for dynamic target

chem_kernel = torch.tensor([[0., 1., 0.],
                             [1., -4., 1.],
                             [0., 1., 0.]], device=device).view(1,1,3,3)

```

```

def gaussian_source(center, sigma, strength):
    cx, cy = center
    r2 = (XS - cx)**2 + (YS - cy)**2
    g = torch.exp(-r2 / (2*sigma*sigma)) * strength
    return g

def sample_chem_sources(params: ChemParams):
    sources = []
    for _ in range(params.n_sources):
        cx = random.uniform(-0.4, 0.4)
        cy = random.uniform(-0.4, 0.4)
        sigma = random.uniform(params.sigma_min, params.sigma_max)
        strength = random.uniform(params.strength_min, params.strength_max)
        sources.append((cx, cy, sigma, strength))
    return sources

def chem_init_field(B, H, W, params: ChemParams):
    S = sample_chem_sources(params)
    base = sum(gaussian_source((cx,cy), sigma, strength) for
    ↪(cx,cy,sigma,strength) in S)
    base = base.clamp(0, 3.0) / 3.0
    field = base.unsqueeze(0).repeat(B, 1, 1) # (B,H,W)
    return field, S

def chem_step(field, params: ChemParams, sources):
    # diffusion Laplacian
    lap = F.conv2d(field.unsqueeze(1), chem_kernel, padding=1).squeeze(1)
    field = field + params.D * lap - params.decay * field
    # inject sources each step (persistent emitters)
    add = 0
    for (cx,cy,sigma,strength) in sources:
        add = add + gaussian_source((cx,cy), sigma, strength)
    field = (field + 0.02 * add).clamp(0, 1) # mild injection
    return field

```

[18]: *#@title (b) Training stage C: chemical-steered morphogenesis*

```

def rollout_training_chem(model: LiquidMorphCA,
                          cfg: Config,
                          optimizer: torch.optim.Optimizer,
                          iters: int = 600,
                          chem_params: ChemParams = ChemParams()):
    loss_log = []
    for it in range(1, iters+1):
        B, H, W = cfg.BATCH, cfg.H, cfg.W
        dummy_instr = torch.zeros((B, cfg.C_INSTR), device=device)

```



```

state, centers = init_state(B, cfg, dummy_instr)

# ignore shape instruction for chem; use neutral [0,0,1] (arbitrary)
instr = torch.zeros((B, cfg.C_INSTR), device=device); instr[:,2] = 1.0
instr_grid = instr.view(B, cfg.C_INSTR, 1, 1).expand(B, cfg.C_INSTR, H, W)
↪W)

state[:, cfg.C_VISIBLE + cfg.C_HIDDEN : cfg.C_VISIBLE + cfg.C_HIDDEN + ↪
↪cfg.C_INSTR] = instr_grid

# init chemical field in SPARE
chem_field, sources = chem_init_field(B, H, W, chem_params)
state[:, -1] = chem_field

T = random.randint(cfg.ROLLOUT_T_MIN, cfg.ROLLOUT_T_MAX)
mid_t = int(0.75 * T); state_mid = None

for t in range(T):
    # update chemical then do a CA step
    chem_field = chem_step(chem_field, chem_params, sources)
    state[:, -1] = chem_field
    state = model.step(state, steps=1)
    if t == mid_t: state_mid = state

# dynamic target = "cover where chem > thresh"
target_mask = (chem_field > chem_params.threshold).float()
targets = target_mask.unsqueeze(1).repeat(1,3,1,1)

# use existing loss function
loss_end, stats_end = compute_losses(state, targets, cfg)
if state_mid is not None:
    loss_mid, _ = compute_losses(state_mid, targets, cfg)
    total = (1 - cfg.MID_LOSS_W) * loss_end + cfg.MID_LOSS_W * loss_mid
    stats_end["loss_mid"] = float(loss_mid.detach())
else:
    total = loss_end

optimizer.zero_grad(set_to_none=True)
total.backward()
torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
optimizer.step()

loss_log.append(float(total.detach()))
if it % 100 == 0 or it == 1:
    print(f"[Chem] it={it}/{iters} | loss={float(total.detach()):.4f} | ↪
↪"

    f"loss_rgb={stats_end['loss_rgb']:.4f} | ↪
↪alive_sup={stats_end['loss_alive_sup']:.4f}")

```

```
return loss_log
```

```
[19]: #@title (b) Visualization: chemical-steered growth

@torch.no_grad()
def render_growth_chem(model: LiquidMorphCA,
                      cfg: Config,
                      steps: int = 64,
                      chem_params: ChemParams = ChemParams(),
                      seed=None):
    if seed is not None: set_seed(seed)
    B, H, W = 1, cfg.H, cfg.W
    instr = torch.tensor([[0,0,1]], dtype=torch.float32, device=device)
    state, centers = init_state(B, cfg, instr)

    chem_field, sources = chem_init_field(B, H, W, chem_params)
    state[:, -1] = chem_field

    frames, chems = [], []
    for t in range(steps):
        chem_field = chem_step(chem_field, chem_params, sources)
        state[:, -1] = chem_field
        state = model.step_no_grad(state, steps=1)
        frames.append(torch.sigmoid(state[0,:3]).permute(1,2,0).cpu().numpy())
        chems.append(chem_field[0].cpu().numpy())

    fig, axes = plt.subplots(2, 3, figsize=(10,6))
    axes[0,0].imshow(frames[0]); axes[0,0].set_title("RGB Step 0");
    ↪ axes[0,0].axis('off')
    axes[0,1].imshow(frames[steps//2]); axes[0,1].set_title(f"RGB {steps//2}");
    ↪ axes[0,1].axis('off')
    axes[0,2].imshow(frames[-1]); axes[0,2].set_title(f"RGB {steps-1}");
    ↪ axes[0,2].axis('off')
    axes[1,0].imshow(chems[0], cmap='gray'); axes[1,0].set_title("Chem_
    ↪ 0"); axes[1,0].axis('off')
    axes[1,1].imshow(chems[steps//2], cmap='gray'); axes[1,1].set_title(f"Chem_
    ↪ {steps//2}"); axes[1,1].axis('off')
    axes[1,2].imshow(chems[-1], cmap='gray'); axes[1,2].set_title(f"Chem_
    ↪ {steps-1}"); axes[1,2].axis('off')
    plt.show()

# Example:
# _ = rollout_training_chem(model, cfg, optimizer, iters=300,
    ↪ chem_params=ChemParams())
# render_growth_chem(model, cfg, steps=64, chem_params=ChemParams(n_sources=3),
    ↪ seed=42)
```

[20]: *#@title (c) Morphological descriptors & novelty score*

```
def morph_descriptors_from_state(state):  
    """  
    Extract a compact descriptor from the final alive mask.  
    Returns vector  $d$  in  $\mathbb{R}^6$ : [area, perimeter, circularity, xc, yc, anisotropy]  
    """  
    with torch.no_grad():  
        alive = torch.sigmoid(state[:, 3:4]) # (B,1,H,W)  
        B, _, H, W = alive.shape  
        # binarize at 0.5  
        mask = (alive > 0.5).float()  
  
        # area  
        area = mask.mean(dim=(2,3)) # (B,1)  
  
        # perimeter via Sobel magnitude  
        kx = torch.tensor([[-1,0,1],[-2,0,2],[-1,0,1]], device=device,   
dtype=torch.float32).view(1,1,3,3)  
        ky = torch.tensor([[-1,-2,-1],[0,0,0],[1,2,1]], device=device,   
dtype=torch.float32).view(1,1,3,3)  
        gx = F.conv2d(mask, kx, padding=1); gy = F.conv2d(mask, ky, padding=1)  
        perim = (gx.abs() + gy.abs()).mean(dim=(2,3)) # (B,1)  
  
        # circularity:  $4A / P^2$  (avoid div by zero)  
        circ = (4*math.pi*area) / (perim*perim + 1e-6)  
  
        # centroid & anisotropy from second moments  
        ys = torch.linspace(0, 1, H, device=device).view(1,1,H,1)  
        xs = torch.linspace(0, 1, W, device=device).view(1,1,1,W)  
        m = mask + 1e-6  
        mx = (m*xs).sum(dim=(2,3)) / m.sum(dim=(2,3))  
        my = (m*ys).sum(dim=(2,3)) / m.sum(dim=(2,3))  
        # covariance  
        cx = ((xs - mx.view(B,1,1,1))**2 * m).sum(dim=(2,3)) / m.sum(dim=(2,3))  
        cy = ((ys - my.view(B,1,1,1))**2 * m).sum(dim=(2,3)) / m.sum(dim=(2,3))  
        anis = torch.maximum(cx, cy) / (torch.minimum(cx, cy) + 1e-6)  
  
        d = torch.cat([area, perim, circ, mx, my, anis], dim=1) # (B,6)  
        return d  
  
def novelty_score(d, archive, k=5):  
    """  
    Mean Euclidean distance to k nearest in archive (torch tensors).  
    """  
    if len(archive) == 0:  
        return torch.tensor([10.0]*d.shape[0], device=d.device)
```

```

D = torch.stack(archive, dim=0) # (N,6)
# pairwise distance: ||d - D||_2
# expand dims for broadcasting
diff = d.unsqueeze(1) - D.unsqueeze(0) # (B,N,6)
dist = torch.sqrt((diff*diff).sum(dim=2) + 1e-8) # (B,N)
k = min(k, dist.shape[1])
topk = torch.topk(dist, k=k, largest=False).values # (B,k)
return topk.mean(dim=1) # (B,)

```

```

<>:6: SyntaxWarning: invalid escape sequence '\i'
<>:6: SyntaxWarning: invalid escape sequence '\i'
/tmp/ipython-input-2387349102.py:6: SyntaxWarning: invalid escape sequence '\i'
Returns vector d \in R^6: [area, perimeter, circularity, xc, yc, anisotropy]

```

```

[26]: #@title (c) ASAL-style search over environments (novelty archive)

from collections import deque

def sample_random_env():
    p = ChemParams()
    p.D = random.uniform(0.15, 0.35)
    p.decay = random.uniform(0.005, 0.03)
    p.n_sources = random.randint(1, 3)
    p.sigma_min = 0.06; p.sigma_max = 0.22
    p.strength_min = 0.8; p.strength_max = 1.4
    p.threshold = random.uniform(0.25, 0.45)
    return p

@torch.no_grad()
def simulate_env_once(model, cfg, chem_params: ChemParams, steps: int = 64):
    B, H, W = 1, cfg.H, cfg.W
    instr = torch.tensor([[0,0,1]], dtype=torch.float32, device=device)
    state, centers = init_state(B, cfg, instr)
    chem_field, sources = chem_init_field(B, H, W, chem_params)
    state[:, -1] = chem_field
    for t in range(steps):
        chem_field = chem_step(chem_field, chem_params, sources)
        state[:, -1] = chem_field
        state = model.step_no_grad(state, steps=1)
    desc = morph_descriptors_from_state(state)[0] # (6,)
    rgb = torch.sigmoid(state[0,:3]).permute(1,2,0).cpu().numpy()
    return desc, rgb, chem_params

def asal_search(model, cfg, num_envs: int = 40, steps: int = 64, top_k: int = 8):
    archive = [] # descriptors (torch tensors)

```

```

    hall: List[Tuple[np.ndarray, ChemParams, torch.Tensor]] = [] # (rgb,
↳ params, desc)
    for i in range(1, num_envs+1):
        params = sample_random_env()
        d, img, env = simulate_env_once(model, cfg, params, steps=steps)
        d = d.to(device)
        nov = float(novelty_score(d.unsqueeze(0), archive, k=5)[0].detach())
        archive.append(d)
        hall.append((img, env, d.cpu()))
        if i % 10 == 0 or i == num_envs:
            print(f"[ASAL] env {i}/{num_envs} | novelty={nov:.3f} |
↳ D={float(params.D):.3f} "
                  f"decay={float(params.decay):.3f} src={params.n_sources}
↳ thr={params.threshold:.2f}")

        # pick top_k most novel (by novelty w.r.t. the rest)
        # recompute novelty against all for fair ranking
        Ds = torch.stack([d for (_,_,d) in hall], dim=0).to(device)
        nov_all = novelty_score(Ds, [desc.to(device) for (_,_, desc) in hall],
↳ k=5).cpu().numpy()
        idxs = np.argsort(-nov_all)[:top_k]

        # visualize
        ncols = top_k
        plt.figure(figsize=(2.2*ncols, 2.2))
        for j, idx in enumerate(idxs):
            plt.subplot(1, ncols, j+1)
            plt.imshow(hall[idx][0]); plt.axis('off')
            p = hall[idx][1]
            plt.title(f"D={p.D:.2f}\ndec={p.decay:.2f}\nsrc={p.n_sources}\nth={p.
↳ threshold:.2f}", fontsize=8)
            plt.suptitle("ASAL: most novel emergent morphologies", y=1.02)
            plt.tight_layout()
            plt.show()

    return [hall[i] for i in idxs] # winners: [(rgb, params, desc), ...]

```

[22]: *#@title (c) Optional: Add ASAL winners to curriculum (fine-tune)*

```

def finetune_on_asal_winners(model, cfg, optimizer, winners, iters_per_env: int,
↳ 150, steps: int = 64):
    for wi, (_, params, _) in enumerate(winners, 1):
        print(f"[ASAL-train] env {wi}/{len(winners)}")
        chem_params = params
        _ = rollout_training_chem(model, cfg, optimizer, iters=iters_per_env,
↳ chem_params=chem_params)

```

```
print("ASAL curriculum fine-tune complete.")
```

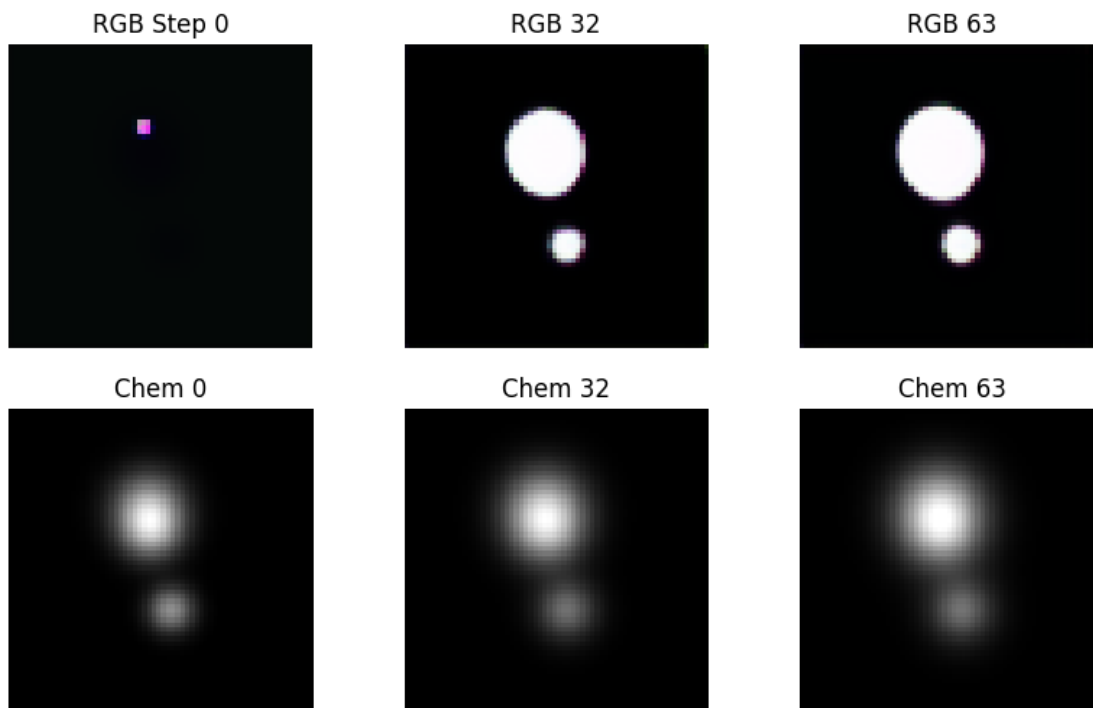
```
[23]: # 1) (a) Generalization: quick fine-tune + visualize
_ = rollout_training_param(model, cfg, optimizer, iters=400)
render_growth_param(model, cfg, shape_id=1, scale=1.2, theta=0.3, steps=64,
↳seed=1)
```

```
[A-gen] it=1/400 | loss=4.7074 | loss_rgb=0.8533 | alive_sup=5.5192
[A-gen] it=100/400 | loss=0.0735 | loss_rgb=0.0233 | alive_sup=0.0684
[A-gen] it=200/400 | loss=0.0647 | loss_rgb=0.0199 | alive_sup=0.0614
[A-gen] it=300/400 | loss=0.0460 | loss_rgb=0.0144 | alive_sup=0.0410
[A-gen] it=400/400 | loss=0.0370 | loss_rgb=0.0121 | alive_sup=0.0317
```



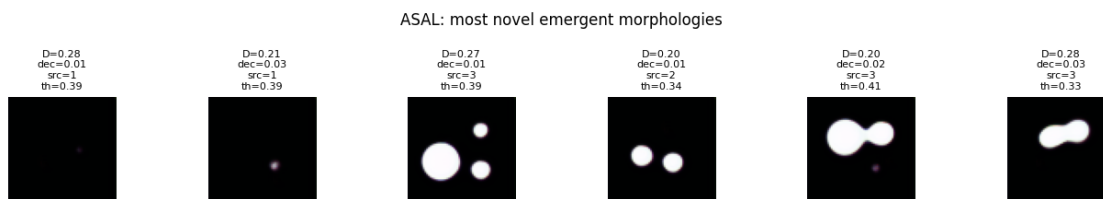
```
[24]: # 2) (b) Chemical steering: fine-tune + visualize
_ = rollout_training_chem(model, cfg, optimizer, iters=300,
↳chem_params=ChemParams(n_sources=2))
render_growth_chem(model, cfg, steps=64, chem_params=ChemParams(n_sources=3,
↳threshold=0.3), seed=7)
```

```
[Chem] it=1/300 | loss=4.9562 | loss_rgb=0.9858 | alive_sup=5.5302
[Chem] it=100/300 | loss=0.1180 | loss_rgb=0.1010 | alive_sup=0.0254
[Chem] it=200/300 | loss=0.0220 | loss_rgb=0.0079 | alive_sup=0.0205
[Chem] it=300/300 | loss=0.0216 | loss_rgb=0.0045 | alive_sup=0.0115
```



```
[27]: # 3) (c) ASAL search + optional curriculum
winners = asal_search(model, cfg, num_envs=40, steps=64, top_k=6)
finetune_on_asal_winners(model, cfg, optimizer, winners, iters_per_env=120,
↳ steps=64)
```

```
[ASAL] env 10/40 | novelty=11.254 | D=0.206 decay=0.029 src=1 thr=0.39
[ASAL] env 20/40 | novelty=0.594 | D=0.166 decay=0.009 src=1 thr=0.29
[ASAL] env 30/40 | novelty=3.218 | D=0.336 decay=0.014 src=1 thr=0.34
[ASAL] env 40/40 | novelty=5.493 | D=0.200 decay=0.015 src=2 thr=0.34
```



```
[ASAL-train] env 1/6
[Chem] it=1/120 | loss=0.0138 | loss_rgb=0.0044 | alive_sup=0.0130
[Chem] it=100/120 | loss=0.0102 | loss_rgb=0.0035 | alive_sup=0.0096
[ASAL-train] env 2/6
[Chem] it=1/120 | loss=0.0141 | loss_rgb=0.0049 | alive_sup=0.0124
[Chem] it=100/120 | loss=0.0052 | loss_rgb=0.0015 | alive_sup=0.0046
```

```
[ASAL-train] env 3/6
[Chem] it=1/120 | loss=0.2237 | loss_rgb=0.0385 | alive_sup=0.0683
[Chem] it=100/120 | loss=0.0315 | loss_rgb=0.0070 | alive_sup=0.0222
[ASAL-train] env 4/6
[Chem] it=1/120 | loss=0.1051 | loss_rgb=0.0296 | alive_sup=0.0970
[Chem] it=100/120 | loss=0.0094 | loss_rgb=0.0032 | alive_sup=0.0083
[ASAL-train] env 5/6
[Chem] it=1/120 | loss=0.1064 | loss_rgb=0.0287 | alive_sup=0.1214
[Chem] it=100/120 | loss=0.0086 | loss_rgb=0.0027 | alive_sup=0.0072
[ASAL-train] env 6/6
[Chem] it=1/120 | loss=0.8061 | loss_rgb=0.1849 | alive_sup=0.8596
[Chem] it=100/120 | loss=0.0397 | loss_rgb=0.0143 | alive_sup=0.0375
ASAL curriculum fine-tune complete.
```

[]:

This notebook was converted with convert.ploomber.io